

# LUKASZ ZAJAC LAB 1 ADZD

## ASSIGNMENT:

Korzystając z operacji na RDD i pliku z logami Apache2:

1. Narysować wykres rankingu najpopularniejszych stron źródłowych (referer), z wyłączeniem domeny semicomplete.com
2. Narysować wykres rankingu pobieranych plików graficznych
3. Narysować średni dobowy wykres liczby wywołań (request) na minutę w zależności od czasu.
4. Narysować wykres (pie chart) popularności systemów operacyjnych klienta
5. Oszacować jaki procent ruchu generują boty. Porównać co najmniej 2 sposoby oszacowania.
6. Narysować mapę z lokalizacjami klientów

## ✓ Introduction to Apache Spark Resilient Distributed Datasets (RDD)

```
!apt-get install openjdk-8-jdk-headless -qq > /dev/null
!wget -q https://archive.apache.org/dist/spark/spark-3.1.2/spark-3.1.2-bin-hadoop3.2.tgz
!tar -xvf spark-3.1.2-bin-hadoop3.2.tgz
!pip install -q findspark
import os
os.environ["JAVA_HOME"] = "/usr/lib/jvm/java-8-openjdk-amd64"
os.environ["SPARK_HOME"] = "/content/spark-3.1.2-bin-hadoop3.2"
import findspark
findspark.init()
from pyspark.sql import SparkSession

spark = SparkSession.builder.master("local[*]").getOrCreate()
```

## ✓ Initialization

Starting Spark session:

```
from pyspark.sql import SparkSession
spark = SparkSession.builder.getOrCreate()
```

spark

 **SparkSession - in-memory**

**SparkContext**

[Spark UI](#)

Version

v3.1.2

Master

local[\*]

AppName

pyspark-shell

```
sc = spark.sparkContext
```

## ✓ Reading text files

Download a file locally (just for the demo, for larger files use HDFS or S3):

```
import urllib.request
```

```
urllib.request.urlretrieve("http://raw.githubusercontent.com/elastic/examples/master/Common%20Data%20Formats/apache_logs/apache_logs",'')
```

## Create a first RDD

```
lines = sc.textFile("apache_logs")
```

Perform a *transformation*

```
lineLengths = lines.map(lambda s: len(s))
```

Perform an *action*

```
totalLength = lineLengths.reduce(lambda a, b: a + b)
```

Finally evaluate (laziness!)

totalLength

 2360789

- Custom functions

```
def myFunc(s):
    words = s.split(" ")
    return len(words)
```

```
wordCounts = lines.map(myFunc)
```

```
wordCounts.max()
```

→ 43

```
wordCounts.histogram([0,10,20,30,40,50])
```

$$\Rightarrow ([0, 10, 20, 30, 40, 50], [0, 4901, 4719, 378, 2])$$

```
wordCounts.count()
```

 10000

- Map-Reduce pattern

Let's look closer at our log file format:

```
99.158.0.150 - - [18/May/2015:04:05:37 +0000] "GET /images/jordan-80.png HTTP/1.1" 200 6146 "http://www.semicomplete.com/articles/dynamic-dns-wi
```

### Count lines containing "Chrome":

```
chromeLines = lines.map(lambda s: (1 if "Chrome" in s else 0))
```

```
counts = chromeLines.reduce(lambda a, b: a + b)
```

counts

3172

- Key-Value Pairs

Count IP addresses

```
regex = '([(\d\.)]+)'

import re

ip = re.match(regex, '99.158.0.150 - - ').group(0)

print(ip)

↵ 99.158.0.150

ips = lines.map(lambda s: (re.match(regex, s).group(0) ,1))

ipCounts = ips.reduceByKey(lambda a, b: a + b)
```

Get results

```
ipCounts.collect()

ipCounts.take(20)

↵ [('83.149.9.216', 23),
   ('110.136.166.128', 6),
   ('46.105.14.53', 364),
   ('123.125.71.35', 1),
   ('200.49.190.100', 1),
   ('67.214.178.190', 2),
   ('207.241.237.220', 18),
   ('207.241.237.227', 17),
   ('207.241.237.228', 16),
   ('218.30.103.62', 16),
   ('71.212.224.97', 5),
   ('86.1.76.62', 6),
   ('107.170.41.69', 10),
   ('208.115.111.72', 83),
   ('174.37.205.76', 1),
   ('134.76.249.10', 8),
   ('123.125.71.117', 1),
   ('65.19.138.34', 3),
   ('5.102.173.71', 2),
   ('187.45.193.158', 8)]
```

Sort by values (ascending):

```
ipCounts.takeOrdered(10, lambda pair: pair[1])

↵ [('123.125.71.35', 1),
   ('200.49.190.100', 1),
   ('174.37.205.76', 1),
   ('123.125.71.117', 1),
   ('180.76.5.27', 1),
   ('93.164.60.142', 1),
   ('94.175.135.116', 1),
   ('172.56.29.98', 1),
   ('173.252.73.114', 1),
   ('54.226.176.99', 1)]
```

Sort by values (descending):

```
topIPs = ipCounts.takeOrdered(10, lambda pair: -pair[1])
topIPs

↵ [('66.249.73.135', 482),
   ('46.105.14.53', 364),
   ('130.237.218.86', 357),
   ('75.97.9.59', 273),
   ('50.16.19.13', 113),
   ('209.85.238.199', 102),
   ('68.180.224.225', 99),
   ('100.43.83.137', 84),
```

```
( '208.115.111.72', 83),
( '198.46.149.143', 82)]
```

## Plotting

```
import pandas as pd
```

```
df = pd.DataFrame(topIPs, columns=['IP', 'Count'])
df
```


|   | IP             | Count |
|---|----------------|-------|
| 0 | 66.249.73.135  | 482   |
| 1 | 46.105.14.53   | 364   |
| 2 | 130.237.218.86 | 357   |
| 3 | 75.97.9.59     | 273   |
| 4 | 50.16.19.13    | 113   |
| 5 | 209.85.238.199 | 102   |
| 6 | 68.180.224.225 | 99    |
| 7 | 100.43.83.137  | 84    |
| 8 | 208.115.111.72 | 83    |
| 9 | 198.46.149.143 | 82    |

Next steps:

[Generate code with df](#)
[View recommended plots](#)
[New interactive sheet](#)

```
import matplotlib.pyplot as plt
```

```
df.plot.barh(x = 'IP')
```



## Closures and accumulators

Local variables are not passed to closures or functions, they remain local!

```
counter = 0
```

```
# Wrong: Don't do this!!
def increment_counter(x):
    global counter
    counter += x
```

```
wordCounts.foreach(increment_counter)

print("Counter value: ", counter)

↗ Counter value: 0
```

Use *accumulators* instead

```
accum = sc.accumulator(0)
accum

↗ Accumulator<id=0, value=0>

# Correct way is to use accumulators
def increment_counter(x):
    accum.add(x)

wordCounts.foreach(increment_counter)

print("Counter value: ", accum)

↗ Counter value: 197956
```

## ✓ Assignment solutions

Narysować wykres rankingu najpopularniejszych stron źródłowych (referer), z wyłączeniem domeny semicomplete.com

```
lines.collect()

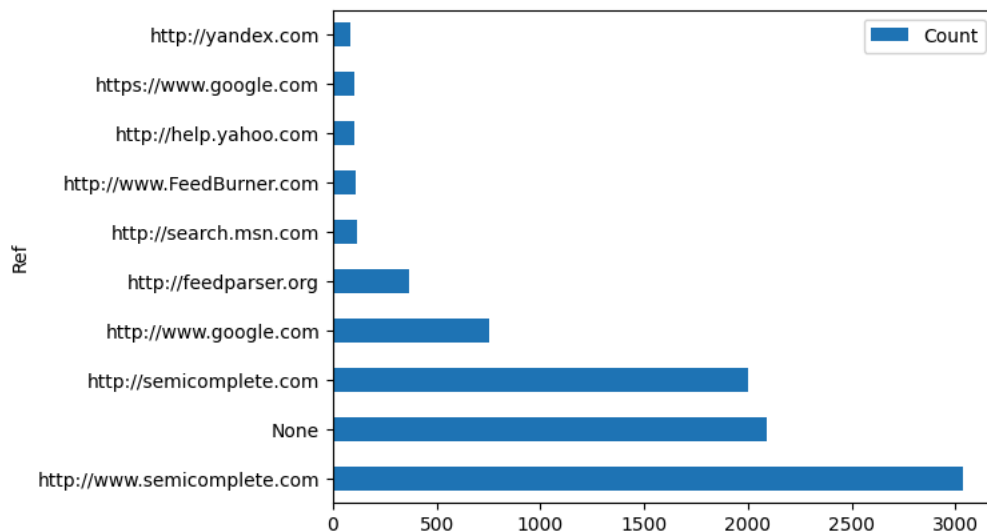
import regex as re
ref_regex = 'https?:\\/(\\[[^\\]\\"]*)'
test_str = ""
17/May/2015:10:05:03 +0000] "GET /presentations/logstash-monitorama-2013/images/kibana-search.png HTTP/1.1" 200 203023 "http://semicom
'83.149.9.216 - - [17/May/2015:10:05:43 +0000] "GET /presentations/logstash-monitorama-2013/images/kibana-dashboard3.png HTTP/1.1" 200
'83.149.9.216 - - [17/May/2015:10:05:47 +0000] "GET /presentations/logstash-monitorama-2013/plugin/highlight/highlight.js HTTP/1.1" 200
'83.149.9.216 - - [17/May/2015:10:05:12 +0000] "GET /presentations/logstash-monitorama-2013/plugin/zoom-js/zoom.js HTTP/1.1" 200 7697
'83.149.9.216 - - [17/May/2015:10:05:07 +0000] "GET /presentations/logstash-monitorama-2013/plugin/notes/notes.js HTTP/1.1" 200 2892
'83.149.9.216 - - [17/May/2015:10:05:34 +0000] "GET /presentations/logstash-monitorama-2013/images/sad-medic.png HTTP/1.1" 200 430406
""

referers = lines.map( lambda l: (None,1) if re.search(ref_regex,l) is None else (re.search(ref_regex,l).group(0),1) )
referers.collect()
referers_count = referers.reduceByKey(lambda a, b: a + b)
referers_count = referers_count.collect()

df = pd.DataFrame(referers_count, columns=['Ref', 'Count'])
df.sort_values(by=['Count'], ascending=False, inplace=True)

df[:10].plot.barh(x = 'Ref')
```

<Axes: ylabel='Ref'>



Narysować wykres rankingu pobieranych plików graficznych


```
#zakladam ze obrazki maja rozszerzenia png lub jpg
file_regex = r'([^\s]+\.(?:png|jpg|jpeg|gif))'

# Find all matches

image_files = lines.map( lambda l: (None,1) if re.search(file_regex,l) is None else (re.search(file_regex,l).group(0),1) )
image_files.collect()
image_files_count = image_files.reduceByKey(lambda a, b: a + b)
image_files_count = image_files_count.collect()
df = pd.DataFrame(image_files_count, columns=['File', 'Count'])

df.sort_values(by=['Count'], ascending=False, inplace=True)
df[1:11].plot.barh(x = 'File')
```

<Axes: ylabel='File'>



Narysować średni dobowy wykres liczby wywołań (request) na minutę w zależności od czasu.

```
time_regex = r'^[^\s]*\:[0-9]{2}\:[0-9]{2}\:[0-9]{2} \+[0-9]{4}\$'
```


```
time_records = lines.map( lambda l: (None,1) if re.search(time_regex,l) is None else (re.search(time_regex,l).group(1),1) )
time_records.collect()
time_records_count = time_records.reduceByKey(lambda a, b: a + b)
time_records_count = time_records_count.collect()
df = pd.DataFrame(time_records_count, columns=['Time', 'Count'])
```

```

u1 = pd.DataFrame(time_records_count, columns=['Time', 'Count'])
df.sort_values(by=['Time'], ascending=True, inplace=True)
df.plot(x='Time')

```

<Axes: xlabel='Time'>



Narysować wykres (pie chart) popularności systemów operacyjnych klienta

```


#zakladam ze obrazki maja rozszerzenia png lub jpg
os_regex = r'(Windows NT|[L|l]inux|Mac OS)'

# Find all matches

os_records = lines.map( lambda l: ("Unknown",1) if re.search(os_regex,l) is None else (re.search(os_regex,l).group(0),1) )
os_records.collect()
os_records_count = os_records.reduceByKey(lambda a, b: a + b)
os_records_count = os_records_count.collect()
df = pd.DataFrame(os_records_count, columns=['OS', 'Count'])
df.plot.pie(y='Count', labels=df['OS'])

```

<Axes: ylabel='Count'>



Oszacować jaki procent ruchu generują boty. Porównać co najmniej 2 sposoby oszacowania.

```

#zakladam ze obrazki maja rozszerzenia png lub jpg
bot_regex = r'Googlebot'

# Find all matches

```

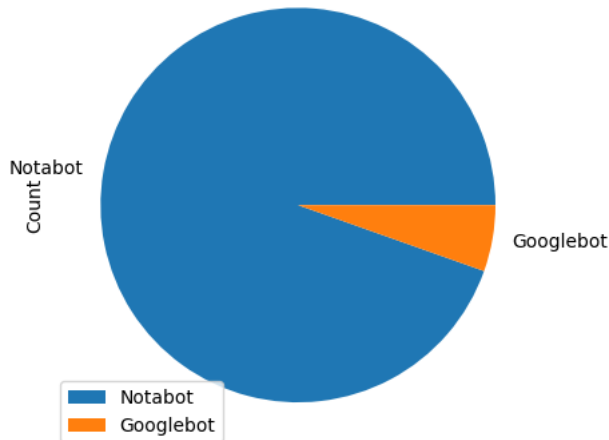
```

bot_regex_found = lines.map( lambda l: ("Notabot",1) if re.search(bot_regex,l) is None else (re.search(bot_regex,l).group(0),1) )
bot_regex_found.collect()
bot_regex_count = bot_regex_found.reduceByKey(lambda a, b: a + b)
bot_regex_count = bot_regex_count.collect()
df = pd.DataFrame(bot_regex_count, columns=['Type', 'Count'])
df.plot.pie(y='Count', labels=df['Type'])

print(df['Count'][0]/df['Count'].sum())

```

0.9457



```

#zakladam ze obrazki maja rozszerzenia png lub jpg
bot_regex = r'Googlebot'

```

```

# Find all matches

```

```

bot_acumm = sc.accumulator(0)
all_acumm = sc.accumulator(0)

```

```

def increment_counter(x):
    if x == "Googlebot":
        bot_acumm.add(1)
        all_acumm.add(1)

```

```

bot_regex_found = lines.map( lambda l: "Notabot" if re.search(bot_regex,l) is None else re.search(bot_regex,l).group(0) )
bot_regex_found.foreach(increment_counter)

```

```

print(1 - bot_acumm.value/all_acumm.value)

```

0.9457

Narysować mapę z lokalizacjami klientów

```

pip install geoip2

```

```

import requests

```

```

def get_ip_geolocation(ip):
    # You can replace 'your_token' with your IP geolocation API token
    response = requests.get(f'https://ipinfo.io/{ip}/geo', params={'token': ADD_YOUR_TOKEN })
    if response.status_code == 200:
        longitude, latitude = response.json()['loc'].split(',')
        return [longitude,latitude]
    else:
        return [0,0]

```

```

get_ip_geolocation("83.149.9.216")

```


['55.7522', '37.6156']




```
#Import Geocoder
import geocoder

def printDetails(ip):
    return get_ip_geolocation(ip)

printDetails("83.149.9.216")
```

 ['55.7522', '37.6156']

Start coding or [generate](#) with AI.

 PythonRDD[84] at collect at <ipython-input-91-4cd22bc3d850>:1



```
ip_loc = ipCounts.map( lambda x: printDetails(x[0]) )

locations = ip_loc.collect()

locations_l = [l for l in locations if l is not None]

locations_df = pd.DataFrame(locations_l, columns=['latitude', 'longitude'])

locations_df
```

|      | latitude | longitude |
|------|----------|-----------|
| 0    | 55.7522  | 37.6156   |
| 1    | -7.7625  | 110.4317  |
| 2    | 48.5839  | 7.7455    |
| 3    | 31.2222  | 121.4581  |
| 4    | 14.6407  | -90.5133  |
| ...  | ...      | ...       |
| 1748 | 50.6942  | 3.1746    |
| 1749 | 47.0056  | 28.8575   |
| 1750 | 32.6469  | -97.3325  |
| 1751 | 59.9386  | 30.3141   |
| 1752 | 25.7743  | -80.1937  |

1753 rows × 2 columns

Next steps:

[Generate code with locations\\_df](#)

 [View recommended plots](#)

[New interactive sheet](#)

```
locations_df['latitude'] = locations_df['latitude'].astype(float)
locations_df['longitude'] = locations_df['longitude'].astype(float)
```

```
import folium
```

```
# Create a world map
world_map = folium.Map(zoom_start=2)
```

```
# Add markers for each geolocation
for location in locations_l:
    folium.Marker(location).add_to(world_map)
```

```
# Save the map as an HTML file
world_map.save('ip_geolocation_map.html')
```

```
display(world_map)
```

